```
placement:
  hosts:
    - clienta.lab.example.com
    - serverc.lab.example.com
    - serverd.lab.example.com
    - servere.lab.example.com
---
service_type: osd ❺
service_id: default_drive_group
placement: ❻
  host_pattern: 'server*'
data_devices:
  paths:
    - /dev/vdb
    - /dev/vdc
    - /dev/vdd
```

❶ The `service_type: host` defines the nodes to add after the `cephadm bootstrap` completes. Host `clienta` will be configured as an admin node.

❷ The Ceph Orchestrator deploys one monitor daemon by default. In the file the `service_type: mon` deploys a `Ceph monitor` daemon in the listed hosts.

❸ The `service_type: mgr` deploys a `Ceph Object Gateway` daemon in the listed hosts.

❹ The `service_type: mgr` deploys a `Ceph Manager` daemon in the listed hosts.

❺ The `service_type: osd` deploys a `ceph-osd` daemon in the listed hosts backed by the `/dev/vdb` device.

❻ Defines where and how to deploy the daemons.

4. As the `root` user on the `serverc` node, run the `cephadm bootstrap` command to create the Ceph cluster. Use the service specification file located at `initial-config-primary-cluster.yaml`

```
[root@serverc ~]# cd /root/ceph
```

```
[root@serverc ceph]# cephadm bootstrap --mon-ip=172.25.250.12 \
--apply-spec=initial-config-primary-c luster.yaml \
--initial-dashboard-password=redhat \
--dashboard-password-noupdate \
--allow-fqdn-hostname \
--registry-url=registry.lab.example.com \
--registry-username=registry \
--registry-password=redhat
...output omitted...
Ceph Dashboard is now available at:

             URL: https://serverc.lab.example.com:8443/
            User: admin
        Password: redhat
```