

```

Applying initial-config-primary-cluster.yaml to cluster
Adding ssh key to clienta.lab.example.com
Adding ssh key to serverd.lab.example.com
Adding ssh key to servere.lab.example.com
Added host 'clienta.lab.example.com' with addr '172.25.250.10'
Added host 'serverc.lab.example.com' with addr '172.25.250.12'
Added host 'serverd.lab.example.com' with addr '172.25.250.13'
Added host 'servere.lab.example.com' with addr '172.25.250.14'
Scheduled mon update...
Scheduled rgw.realm.zone update...
Scheduled mgr update...
Scheduled osd.default_drive_group update...

```

You can access the Ceph CLI with:

```

sudo /usr/sbin/cephadm shell --fsid 8896efec-21ea-11ec-b6fe-52540000fa0c -c /etc/
ceph/ceph.conf -k /etc/ceph/ceph.client.admin.keyring

```

Please consider enabling telemetry to help improve Ceph:

```
ceph telemetry on
```

For more information see:

```
https://docs.ceph.com/docs/pacific/mgr/telemetry/
```

Bootstrap complete.

► 5. Verify the status of the Ceph storage cluster.

5.1. Run the cephadm shell.

```

[root@serverc ~]# cephadm shell
...output omitted...
[ceph: root@serverc /]#

```

5.2. Verify that the cluster status is HEALTH_OK.

```

[ceph: root@serverc /]# ceph status
cluster:
  id:      8896efec-21ea-11ec-b6fe-52540000fa0c
  health: HEALTH_OK

services:
  mon: 4 daemons, quorum serverc.lab.example.com,serverd,servere,clienta (age
10s)
  mgr: serverc.lab.example.com.bypxer(active, since 119s), standbys:
serverd.lflgzj, clienta.hloibd, servere.jhegip
  osd: 9 osds: 9 up (since 55s), 9 in (since 75s)

data:
  pools: 1 pools, 1 pgs

```