

- 3.2. Using the `cephadm` shell, verify that the cluster was successfully deployed. Wait for the cluster to finish deploying and reach the `HEALTH_OK` status.

```
[root@serverc cephadm-ansible]# cephadm shell
...output omitted...
[ceph: root@serverc /]# ceph status
cluster:
  id:      0bbab748-30ee-11ec-abc4-52540000fa0c
  health: HEALTH_OK

services:
  mon: 3 daemons, quorum serverc.lab.example.com,servere,serverd (age 2m)
  mgr: serverc.lab.example.com.blxerd(active, since 3m), standbys:
serverd.nibyts, servere.rkpsii
  osd: 9 osds: 9 up (since 2m), 9 in (since 2m)

data:
  pools:   1 pools, 1 pgs
  objects: 0 objects, 0 B
  usage:   46 MiB used, 90 GiB / 90 GiB avail
  pgs:     1 active+clean
```

4. Expand the cluster by adding OSDs to `serverc`, `serverd`, and `servere`. Use the following service specification file.

```
service_type: osd
service_id: default_drive_group
placement:
  hosts:
    - serverc.lab.example.com
    - serverd.lab.example.com
    - servere.lab.example.com
data_devices:
  paths:
    - /dev/vde
    - /dev/vdf
```

- 4.1. Create a service specification file called `osd-spec.yaml`.

```
[ceph: root@serverc /]# cat /tmp/osd-spec.yaml
service_type: osd
service_id: default_drive_group
placement:
  hosts:
    - serverc.lab.example.com
    - serverd.lab.example.com
    - servere.lab.example.com
data_devices:
  paths:
    - /dev/vde
    - /dev/vdf
```

- 4.2. Use the `ceph orch apply` command to add the OSDs to the cluster OSD nodes.