

```
[ceph: root@clienta /]# ceph osd crush tree
ID CLASS WEIGHT TYPE NAME
-1      0.08817 root default
-3      0.02939 host serverc
 0 hdd 0.00980 osd.0
 2 hdd 0.00980 osd.2
 1 ssd 0.00980 osd.1
-5      0.02939 host serverd
 3 hdd 0.00980 osd.3
 7 hdd 0.00980 osd.7
 5 ssd 0.00980 osd.5
-7      0.02939 host servere
 4 hdd 0.00980 osd.4
 8 hdd 0.00980 osd.8
 6 ssd 0.00980 osd.6
```

- 2.3. Add a new CRUSH map rule called `onssd` to target the OSDs with SSD devices.

```
[ceph: root@clienta /]# ceph osd crush rule create-replicated onssd \
default host ssd
```

- 2.4. Use the `ceph osd crush rule ls` command to verify the successful creation of the new rule.

```
[ceph: root@clienta /]# ceph osd crush rule ls
replicated_rule
onssd
```

- 2.5. Create a new replicated pool called `myfast` with 32 placement groups that uses the `onssd` CRUSH map rule.

```
[ceph: root@clienta /]# ceph osd pool create myfast 32 32 onssd
pool 'myfast' created
```

- 2.6. Verify that the placement groups for the pool called `myfast` are only using the OSDs backed by SSD storage. In a previous step, the OSDs are `osd.2`, `osd.5`, and `osd.8`. Retrieve the ID of the pool called `myfast`.

```
[ceph: root@clienta /]# ceph osd lspools
...output omitted...
6 myfast
```

- 2.7. Use the `ceph pg dump pgs_brief` command to list all the PGs in the cluster. The pool ID is the first number in a PG ID. For example, the PG `6.1b` belongs to the pool whose ID is 6.

```
[ceph: root@clienta /]# ceph pg dump pgs_brief
PG_STAT STATE UP UP_PRIMARY ACTING ACTING_PRIMARY
6.1b active+clean [6,5,1] 6 [6,5,1] 6
4.19 active+clean [6,2,5] 6 [6,2,5] 6
```