```
2.1f      active+clean   [0,3,8]         0   [0,3,8]          0
3.1e      active+clean   [2,6,3]         2   [2,6,3]          2
6.1a      active+clean   [6,1,5]         6   [6,1,5]          6
4.18      active+clean   [3,2,6]         3   [3,2,6]          3
2.1e      active+clean   [2,6,5]         2   [2,6,5]          2
3.1f      active+clean   [0,3,4]         0   [0,3,4]          0
6.19      active+clean   [1,5,6]         1   [1,5,6]          1
4.1b      active+clean   [3,2,8]         3   [3,2,8]          3
2.1d      active+clean   [6,7,0]         6   [6,7,0]          6
...output omitted...
```

The pool called `myfast`, whose ID is 6, only uses `osd.1`, `osd.5`, and `osd.6`. These are the only OSDs with SSD drives.

▶ **3.** Create a new CRUSH hierarchy under `root=default-cl260` that has three rack buckets (`rack1`, `rack2`, and `rack3`), each of which contains one host bucket (`hostc`, `hostd`, and `hoste`).

   3.1.   Create a new CRUSH map hierarchy that matches this infrastructure:

```
default-cl260     (root bucket)
    rack1        (rack bucket)
        hostc       (host bucket)
            osd.1
            osd.5
            osd.6

    rack2        (rack bucket)
        hostd       (host bucket)
            osd.0
            osd.3
            osd.4

    rack3        (rack bucket)
        hoste       (host bucket)
            osd.2
            osd.7
            osd.8
```

You should place the three SSDs (in this example are OSDs 1, 5, and 6) on `hostc`. Because in your cluster OSD numbers can be differet, modify the CRUSH map hierarchy accordingly to this requirement.

First, create the buckets with the `ceph osd crush add-bucket` command.

```
[ceph: root@clienta /]# ceph osd crush add-bucket default-cl260 root
added bucket default-cl260 type root to crush map
[ceph: root@clienta /]# ceph osd crush add-bucket rack1 rack
added bucket rack1 type rack to crush map
[ceph: root@clienta /]# ceph osd crush add-bucket hostc host
added bucket hostc type host to crush map
[ceph: root@clienta /]# ceph osd crush add-bucket rack2 rack
added bucket rack2 type rack to crush map
[ceph: root@clienta /]# ceph osd crush add-bucket hostd host
added bucket hostd type host to crush map
```