```
pool 1 'device_health_metrics' replicated size 3 min_size 2 crush_rule 0
 object_hash rjenkins pg_num 1 pgp_num 1 autoscale_mode on last_change 475 flags
 hashpspool stripe_width 0 pg_num_min 1 application mgr_devicehealth
...output omitted...
osd.0 up   in  weight 1 up_from 471 up_thru 471 down_at 470 last_clean_interval
 [457,466) [v2:172.25.250.12:6801/1228351148,v1:172.25.250.12:6802/1228351148]
 [v2:172.25.249.12:6803/1228351148,v1:172.25.249.12:6804/1228351148] exists,up
 cfe311b0-dea9-4c0c-a1ea-42aaac4cb160
...output omitted...
```

3. Set the `full_ratio` and `nearfull_ratio` and verify the values.

3.1. Set the `full_ratio` parameter to 0.97 (97%) and `nearfull_ratio` to 0.9 (90%).

```
[ceph: root@clienta /]# ceph osd set-full-ratio 0.97
osd set-full-ratio 0.97
[ceph: root@clienta /]# ceph osd set-nearfull-ratio 0.9
osd set-nearfull-ratio 0.9
```

3.2. Verify the `full_ratio` and `nearfull_ratio` values. Compare this epoch value with the value from the previous dump of the OSD map. The epoch has incremented two versions because each `ceph osd set-*-ratio` command produces a new OSD map version.

```
[ceph: root@clienta /]# ceph osd dump
epoch 480
fsid 11839bde-156b-11ec-bb71-52540000fa0c
created 2021-09-14T14:50:39.401260+0000
modified 2021-09-27T12:27:38.328351+0000
flags sortbitwise,recovery_deletes,purged_snapdirs,pglog_hardlimit
crush_version 69
full_ratio 0.97
backfillfull_ratio 0.9
nearfull_ratio 0.9
...output omitted...
```

4. Extract and view the OSD map.

4.1. Instead of using the `ceph osd dump` command, use the `ceph osd getmap` command to extract a copy of the OSD map to a binary file and the `osdmaptool` command to view the file.

Use the `ceph osd getmap` command to save a copy of the OSD map in the `map.bin` file.

```
[ceph: root@clienta /]# ceph osd getmap -o map.bin
got osdmap epoch 480
```

4.2. Use the `osdmaptool --print` command to display the text version of the binary OSD map. The output is similar to the output of the `ceph osd dump` command.