

```
[root@clientb ~]# rbd showmapped
id pool namespace image snap device
0 test_pool test - /dev/rbd0
```

- 6. Verify that you can use the RBD image mapped on the `clientb` node like a regular disk block device.

6.1. Format the device with an XFS file system.

```
[root@clientb ~]# mkfs.xfs /dev/rbd0
meta-data=/dev/rbd0             isize=512    agcount=8, agsize=4096 blks
                =               sectsz=512   attr=2, projid32bit=1
                =               crc=1        finobt=1, sparse=1, rmapbt=0
                =               reflink=1
data                =           bsize=4096   blocks=32768, imaxpct=25
                =           sunit=16      swidth=16 blks
naming              =version 2           bsize=4096   ascii-ci=0, ftype=1
log                 =internal log        bsize=4096   blocks=1872, version=2
                =           sectsz=512   sunit=16 blks, lazy-count=1
realtime            =none                extsz=4096   blocks=0, rtextents=0
Discarding blocks...Done.
```

6.2. Create a mount point for the file system.

```
[root@clientb ~]# mkdir /mnt/rbd
```

6.3. Mount the file system created on the `/dev/rbd0` device.

```
[root@clientb ~]# mount /dev/rbd0 /mnt/rbd
```

6.4. Change the ownership of the mount point.

```
[root@clientb ~]# chown admin:admin /mnt/rbd
```

6.5. Review the file-system usage.

```
[root@clientb ~]# df /mnt/rbd
Filesystem      1K-blocks  Used Available Use% Mounted on
/dev/rbd0        123584    7940    115644    7% /mnt/rbd
```

6.6. Add some content to the file system.

```
[root@clientb ~]# dd if=/dev/zero of=/mnt/rbd/test1 bs=10M count=1
1+0 records in
1+0 records out
10485760 bytes (10 MB, 10 MiB) copied, 0.00838799 s, 1.3 GB/s
```

```
[root@clientb ~]# ls /mnt/rbd
test1
```