

- 2.1. In the primary cluster, create a pool called `rbd` with 32 placement groups. Enable the `rbd` client application for the Ceph Block Device and make it usable by the RBD feature.

```
[ceph: root@clienta ~]# ceph osd pool create rbd 32
pool 'rbd' created
[ceph: root@clienta ~]# ceph osd pool application enable rbd rbd
enabled application 'rbd' on pool 'rbd'
[ceph: root@clienta ~]# rbd pool init -p rbd
```

- 2.2. In the secondary cluster, create a pool called `rbd` with 32 placement groups. Enable the `rbd` client application for the Ceph Block Device and make it usable by the RBD feature.

```
[ceph: root@serverf ~]# ceph osd pool create rbd 32
pool 'rbd' created
[ceph: root@serverf ~]# ceph osd pool application enable rbd rbd
enabled application 'rbd' on pool 'rbd'
[ceph: root@serverf ~]# rbd pool init -p rbd
```

- ▶ 3. Create an RBD image called `rbd/test` in your primary Ceph cluster. Map it as a block device, format it with an XFS file system, mount it in `/mnt/rbd` directory, copy some data into it, and then unmount it.

- 3.1. Create the RBD image. Exit the `cephadm` shell and switch to the `root` user. Map the image, and then format it with an XFS file system.

```
[ceph: root@clienta ~]# rbd create test --size 128 --pool rbd
[ceph: root@clienta ~]# exit
exit
[admin@clienta ~]$ sudo -i
[root@clienta ~]# rbd map --pool rbd test
/dev/rbd0
[root@clienta ~]# mkfs.xfs /dev/rbd0
...output omitted...
```

- 3.2. Mount `/dev/rbd0` to the `/mnt/rbd` directory and copy a file to it.

```
[root@clienta ~]# mount /dev/rbd0 /mnt/rbd
[root@clienta ~]# mount | grep rbd
/dev/rbd0 on /mnt/rbd type xfs (rw,relatime,seclabel,attr2,inode64,...)
[root@clienta ~]# cp /etc/ceph/ceph.conf /mnt/rbd/file0
[root@clienta ~]# ls /mnt/rbd
file0
```

- 3.3. Unmount your file system to ensure that the system flushes all data to the Ceph cluster.

```
[root@clienta ~]# umount /mnt/rbd
```