```
[ceph: root@serverf /]# exit
exit
[admin@serverf ~]$ sudo -i
[root@serverf ~]# rbd map --pool rbd test
/dev/rbd0
[root@serverf ~]# mount /dev/rbd0 /mnt/rbd
[root@serverf ~]# mount | grep rbd
/dev/rbd0 on /mnt/rbd type xfs (rw,relatime,seclabel,attr2,inode64,...)
[root@serverf ~]# df -h /mnt/rbd
Filesystem      Size  Used Avail Use% Mounted on
/dev/rbd0       121M  7.8M  113M   7% /mnt/rbd
[root@serverf ~]# ls -l /mnt/rbd
total 4
-rw-r--r--. 1 admin users 177 Sep 30 22:02 file0
[root@serverf ~]# cat /mnt/rbd/file0
# minimal ceph.conf for c315020c-21f0-11ec-b6d6-52540000fa0c
[global]
 fsid = c315020c-21f0-11ec-b6d6-52540000fa0c
 mon_host = [v2:172.25.250.12:3300/0,v1:172.25.250.12:6789/0]
```

5.4. Unmount the file system and unmap the RBD image.

```
[root@serverf ~]# umount /mnt/rbd
[root@serverf ~]# rbd unmap /dev/rbd0
```

6. In this part of the exercise, you will create a pair of snapshots of `rbd/test` on your primary cluster and export the changes between those snapshots as an incremental diff image. You will then import the changes from the incremental diff into your copy of the `rbd/test` image on your secondary cluster.

6.1. In the primary cluster, run the `cephadm` shell and create an initial snapshot called `rbd/test@firstsnap`. Calculate the provisioned and actual disk usage of the `rbd/test` image and its associated snapshots.

```
[root@clienta ~]# cephadm shell
...output omitted...
[ceph: root@clienta /]# rbd snap create rbd/test@firstsnap
Creating snap: 100% complete...done.
[ceph: root@clienta /]# rbd du --pool rbd test
NAME            PROVISIONED  USED
test@firstsnap      128 MiB  36 MiB
test                128 MiB  36 MiB
<TOTAL>             128 MiB  72 MiB
[ceph: root@clienta /]# exit
exit
[root@clienta ~]#
```

6.2. In the secondary cluster, run the `cephadm` shell, create an initial snapshot called `rbd/test@firstsnap`. Calculate the provisioned and actual disk usage of the `rbd/test` image and its associated snapshots.