```
[ceph: root@clienta /]# ceph osd pool create rbd260 32
pool 'rbd260' created
[ceph: root@clienta /]# ceph osd pool application enable rbd260 rbd
enabled application 'rbd' on pool 'rbd260'
[ceph: root@clienta /]# rbd pool init -p rbd260
```

1.3. List the `rbd260` pool details to verify your work. The pool ID might be different in your lab environment.

```
[ceph: root@clienta /]# ceph osd pool ls detail | grep rbd260
pool 7 'rbd260' replicated size 3 min_size 2 crush_rule 0 object_hash
 rjenkins pg_num 32 pgp_num 32 autoscale_mode on last_change 203 flags
 hashpspool,selfmanaged_snaps stripe_width 0 application rbd
```

2. Create a 128 MiB RADOS block device image called `prod260` in the `rbd260` pool. Verify your work.

2.1. Create the 128 MiB `prod260` RBD image in the `rbd260` pool.

```
[ceph: root@clienta /]# rbd create prod260 --size 128 --pool rbd260
```

2.2. List the images in the `rbd260` pool to verify the result.

```
[ceph: root@clienta /]# rbd ls rbd260
prod260
```

3. Map the `prod260` RBD image in the `rbd260` pool to a local block device file by using the kernel RBD client. Format the device with an XFS file system. Mount the file system on the `/mnt/prod260` image and copy the `/etc/resolv.conf` file to the root of this new file system. When done, unmount and unmap the device.

3.1. Exit the `cephadm` shell, then switch to the `root` user. Install the `ceph-common` package on the `clienta` node. Map the `prod260` image from the `rbd260` pool using the kernel RBD client.

```
[ceph: root@clienta /]# exit
exit
[admin@clienta ~]# sudo -i
[root@clienta ~]# yum install -y ceph-common
...output omitted...
Complete!
[root@clienta ~]# rbd map --pool rbd260 prod260
/dev/rbd0
```

3.2. Format the `/dev/rbd0` device with an XFS file system and mount the file system on the device. Change the user and group ownership of the root directory of the new file system to `admin`.