```
    pools:   5 pools, 105 pgs
    objects: 190 objects, 5.3 KiB
    usage:   147 MiB used, 90 GiB / 90 GiB avail
    pgs:     105 active+clean
```

Important

Ensure that the monitor daemons displayed in the `services` section match those of your 3-node production cluster plus the client.

1.2. Open another terminal window and log in to `serverf` as the `admin` user and switch to the `root` user. Verify the health of your backup cluster.

```
[student@workstation ~]$ ssh admin@serverf
...output omitted...
[admin@serverf ~]$ sudo -i
[root@serverf ~]# cephadm shell
[ceph: root@clientf /]# ceph status
...output omitted...
  cluster:
    id:     3c67d550-1fd3-11ec-a0d5-52540000fa0f
    health: HEALTH_OK

  services:
    mon: 1 daemons, quorum serverf.lab.example.com (age 18m)
    mgr: serverf.lab.example.com.qfmyuk(active, since 18m)
    osd: 5 osds: 5 up (since 18m), 5 in (since 47h)
    rgw: 1 daemon active (1 hosts, 1 zones)

  data:
    pools:   5 pools, 105 pgs
    objects: 189 objects, 4.9 KiB
    usage:   82 MiB used, 50 GiB / 50 GiB avail
    pgs:     105 active+clean
```

Important

Ensure that the monitor daemon displayed in the `services` section matches that of your single-node backup cluster.

2. Create a pool called `rbd` in the production cluster with 32 placement groups. In the backup cluster, configure a pool to mirror the data from the `rbd` pool in the production cluster to the backup cluster. Pool-mode mirroring *always* mirrors data between two pools that have the same name in both clusters.

2.1. In the production cluster, create a pool called `rbd` with 32 placement groups. Enable the `rbd` client application for the Ceph Block Device and make it usable by the RBD feature.