

- 3. Create the `rgw_service.yaml` file. Configure the service to start two RGW instances in each of the `serverd` and `servere` hosts. The ports of the RGW instances must start from port 8080. Your file should look like this example.

```
[ceph: root@clienta /]# cat rgw_service.yaml
service_type: rgw
service_id: myrealm.myzone
service_name: rgw.myrealm.myzone
placement:
  count: 4
  hosts:
    - serverd.lab.example.com
    - servere.lab.example.com
spec:
  rgw_frontend_port: 8080
```

- 4. Use the Ceph orchestrator to create an RGW service with the `rgw_service.yaml` file. View the cluster and RGW service status. Verify that there are two daemons per host.

4.1. Use Ceph orchestrator to create the RGW service with the `rgw_service.yaml` file.

```
[ceph: root@clienta /]# ceph orch apply -i rgw_service.yaml
Scheduled rgw.myrealm.myzone update...
```

4.2. View the cluster status and find the status of the new RGW service daemons.

```
[ceph: root@clienta /]# ceph status
cluster:
  id:      2ae6d05a-229a-11ec-925e-52540000fa0c
  health: HEALTH_OK

services:
  mon: 4 daemons, quorum serverc.lab.example.com,clienta,serverd,servere (age
4m)
  mgr: serverc.lab.example.com.aiqepd(active, since 10m), standbys:
clienta.nncugs, serverd.klrkci
  osd: 9 osds: 9 up (since 8m), 9 in (since 9m)
  rgw: 4 daemons active (2 hosts, 1 zones)
...output omitted...
```

4.3. Verify that the orchestrator created two running daemons per node.