

```
[admin@clienta ~]$ sudo yum install -y ceph-common
...output omitted...
Complete!
```

- 2.3. Create the /mnt/cephfs-review mount point directory. Mount the new CephFS file system as a kernel client.

```
[admin@clienta ~]$ sudo mkdir /mnt/cephfs-review
[admin@clienta ~]$ sudo mount.ceph serverc.lab.example.com:/ /mnt/cephfs-review \
-o name=admin
```

- 2.4. Change the ownership of the top-level directory of the mounted file system to user and group admin.

```
[admin@clienta ~]$ sudo chown admin:admin /mnt/cephfs-review
```

3. Create a 10 MB test file called cephfs.test1. Verify that the created data is replicated across all three nodes by showing 30 MB in the cephfs_data pool.
 - 3.1. Use the dd command to create one 10 MB file, and then verify that it triples across the OSD nodes.

```
[admin@clienta ~]$ dd if=/dev/zero of=/mnt/cephfs-review/cephfs.test1 \
bs=1M count=10
10+0 records in
10+0 records out
10485760 bytes (10 MB, 10 MiB) copied, 0.0291862 s, 359 MB/s
[admin@clienta ~]$ sudo cephadm shell -- ceph fs status
Inferring fsid ff97a876-1fd2-11ec-8258-52540000fa0c
Inferring config /var/lib/ceph/ff97a876-1fd2-11ec-8258-52540000fa0c/mon.clienta/
config
Using recent ceph image registry.redhat.io/rhceph/rhceph-5-
rhel8@sha256:6306...47ff
mycephfs - 1 clients
=====
RANK STATE MDS ACTIVITY DNS INOS DIRS CAPS
0 active mycephfs.serverc.nsihbi Reqs: 0 /s 11 14 12 2
POOL TYPE USED AVAIL
cephfs_metadata metadata 120k 28.4G
cephfs_data data 30.0M 28.4G
MDS version: ceph version 16.2.0-117.el8cp
(0e34bb74700060ebfaa22d99b7d2cdc037b28a57) pacific (stable)
```

4. Return to workstation as the student user.

```
[admin@clienta ~]$ exit
[student@workstation ~]$
```