```
7   hdd 0.00980  1.00000 10 GiB 1.0 GiB  28 MiB  44 KiB 1024 MiB 9.0 GiB 10.28
 1.00  38     up
8   hdd 0.00980  1.00000 10 GiB 1.0 GiB  28 MiB  44 KiB 1024 MiB 9.0 GiB 10.28
 1.00  47     up
                    TOTAL 90 GiB 9.2 GiB 255 MiB 274 KiB  9.0 GiB  81 GiB 10.28
```

The following table describes each column of the command output:

Output column	Description
ID	The OSD ID.
CLASS	The type of devices that the OSD uses (HDD, SDD, or NVMe).
WEIGHT	The weight of the OSD in the CRUSH map. By default, this is set to the OSD capacity in TB and is changed by using the `ceph osd crush reweight` command. The weight determines how much data CRUSH places onto the OSD relative to other OSDs. For example, two OSDs with the same weight receive roughly the same number of I/O requests and store approximately the same amount of data.
REWEIGHT	Either the default reweight value or the actual value set by the `ceph osd reweight` command. You can reweight an OSD to temporarily override the CRUSH weight.
SIZE	The total OSD storage capacity.
RAW USE	The utilized OSD storage capacity.
DATA	OSD capacity used by user data.
OMAP	The BlueFS storage that is used to store object map (OMAP) data, which are the key-value pairs stored in RocksDB.
META	The total BlueFS space allocated, or the value of the `bluestore_bluefs_min` setting, whichever is larger. This is the internal BlueStore metadata, which is calculated as the total space allocated to BlueFS minus the estimated OMAP data size.
AVAIL	Free space available on the OSD.
%USE	The percentage of storage capacity used on the OSD.
VAR	The variation above or below the average OSD utilization.
PGS	The number of placement groups on the OSD.
STATUS	The status of the OSD.

Use the `ceph osd perf` command to view OSD performance statistics.

```
[ceph: root@node /]# ceph osd perf
```