

```
[ceph: root@node /]# ceph versions
{
  "mon": {
    "ceph version 16.2.0-117.el8cp (0e34bb74700060ebfaa22d99b7d2cdc037b28a57)
    pacific (stable)": 4
  },
  "mgr": {
    "ceph version 16.2.0-117.el8cp (0e34bb74700060ebfaa22d99b7d2cdc037b28a57)
    pacific (stable)": 4
  },
  "osd": {
    "ceph version 16.2.0-117.el8cp (0e34bb74700060ebfaa22d99b7d2cdc037b28a57)
    pacific (stable)": 9
  },
  "mds": {
    "ceph version 16.2.0-117.el8cp (0e34bb74700060ebfaa22d99b7d2cdc037b28a57)
    pacific (stable)": 3
  },
  "rgw": {
    "ceph version 16.2.0-117.el8cp (0e34bb74700060ebfaa22d99b7d2cdc037b28a57)
    pacific (stable)": 2
  },
  "overall": {
    "ceph version 16.2.0-117.el8cp (0e34bb74700060ebfaa22d99b7d2cdc037b28a57)
    pacific (stable)": 22
  }
}
```

Using the Balancer Module

Red Hat Ceph Storage provides a MGR module called **balancer** that automatically optimizes the placement of PGs across OSDs to achieve a balanced distribution. This module can also be run manually.

The **balancer** module does not run if the cluster is not in the **HEALTH_OK** state. When the cluster is healthy, it throttles its changes so that it keeps the number of PGs that need to be moved under a 5% threshold. Configure the **target_max_misplaced_ratio** MGR setting to adjust this threshold:

```
[ceph: root@node /]# ceph config set mgr.* target_max_misplaced_ratio .10
```

The **balancer** module is enabled by default. Use the **ceph balancer on** and **ceph balancer off** commands to enable or disable the balancer.

Use the **ceph balancer status** command to display the balancer status.

```
[ceph: root@node /]# ceph balancer status
```

Automated Balancing

Automated balancing uses one of the following modes: