

```
[ceph: root@clienta /]# ssh admin@serverc sudo journalctl \
-u ceph-ff97a876-1fd2-11ec-8258-52540000fa0c@osd.2.service | grep systemd
admin@serverc's password: redhat
...output omitted...
Sep 30 01:57:36 serverc.lab.example.com systemd[1]: Stopping Ceph osd.2 for
ff97a876-1fd2-11ec-8258-52540000fa0c...
Sep 30 01:57:37 serverc.lab.example.com systemd[1]: ceph-
ff97a876-1fd2-11ec-8258-52540000fa0c@osd.2.service: Succeeded.
Sep 30 01:57:37 serverc.lab.example.com systemd[1]: Stopped Ceph osd.2 for
ff97a876-1fd2-11ec-8258-52540000fa0c.
Sep 30 02:00:12 serverc.lab.example.com systemd[1]: Starting Ceph osd.2 for
ff97a876-1fd2-11ec-8258-52540000fa0c...
Sep 30 02:00:13 serverc.lab.example.com systemd[1]: Started Ceph osd.2 for
ff97a876-1fd2-11ec-8258-52540000fa0c.
```

- 9. Mark the `osd.4` daemon as being out of the cluster and observe how it affects the cluster status. Then, mark the `osd.4` daemon as being in the cluster again.

9.1. Mark the `osd.4` daemon as being out of the cluster. Verify that the `osd.4` daemon is marked out of the cluster and notice that the OSD's weight is now 0.

```
[ceph: root@clienta /]# ceph osd out 4
marked out osd.4.
[ceph: root@clienta /]# ceph osd stat
9 osds: 9 up (since 2m), 8 in (since 3s); epoch: e312
[ceph: root@clienta /]# ceph osd tree
ID CLASS WEIGHT TYPE NAME STATUS REWEIGHT PRI-AFF
-1 0.08817 root default
-3 0.02939 host serverc
0 hdd 0.00980 osd.0 up 1.00000 1.00000
1 hdd 0.00980 osd.1 up 1.00000 1.00000
2 hdd 0.00980 osd.2 up 1.00000 1.00000
-7 0.02939 host serverd
3 hdd 0.00980 osd.3 up 1.00000 1.00000
5 hdd 0.00980 osd.5 up 1.00000 1.00000
7 hdd 0.00980 osd.7 up 1.00000 1.00000
-5 0.02939 host servere
4 hdd 0.00980 osd.4 up 0 1.00000
6 hdd 0.00980 osd.6 up 1.00000 1.00000
8 hdd 0.00980 osd.8 up 1.00000 1.00000
```



Note

Ceph recreates the missing object replicas previously available on the `osd.4` daemon on different OSDs. You can trace the recovery of the objects using the `ceph status` or the `ceph -w` commands.

9.2. Mark the `osd.4` daemon as being in again.

```
[ceph: root@clienta /]# ceph osd in 4
marked in osd.4.
```