

```
[ceph: root@clienta /]# ceph mgr services
{
  "dashboard": "https://172.25.250.12:8443/",
  "prometheus": "http://172.25.250.12:9283/"
}
```

**Note**

Your output might be different depending on which MGR node is active in your lab environment.

2. You receive an alert that an OSD is down. Identify which OSD is down. Identify on which node the down OSD runs, and start the OSD.

- 2.1. Verify cluster health.

```
[ceph: root@clienta /]# ceph health detail
HEALTH_WARN 1 osds down; Degraded data redundancy: 72/666 objects degraded
(10.811%), 14 pgs degraded, 50 pgs undersized
[WRN] OSD_DOWN: 1 osds down
    osd.6 (root=default,host=servere) is down
[WRN] PG_DEGRADED: Degraded data redundancy: 72/666 objects degraded (10.811%), 14
pgs degraded, 50 pgs undersized
    pg 2.0 is stuck undersized for 61s, current state active+undersized, last
    acting [3,0]
    pg 2.1 is stuck undersized for 61s, current state active+undersized, last
    acting [2,3]
    pg 2.6 is stuck undersized for 61s, current state active+undersized, last
    acting [1,3]
    pg 2.7 is stuck undersized for 61s, current state active+undersized, last
    acting [3,2]
...output omitted...
```

- 2.2. Identify which OSD is down.

```
[ceph: root@clienta /]# ceph osd tree | grep -i down
6    hdd    0.00980          osd.6          down    1.00000    1.00000
```

- 2.3. Identify on which host the down OSD runs.

```
[ceph: root@clienta /]# ceph osd find osd.6 | grep host
    "host": "servere.lab.example.com",
    "host": "servere",
```

- 2.4. Start the OSD.

```
[ceph: root@clienta /]# ceph orch daemon start osd.6
Scheduled to start osd.6 on host 'servere.lab.example.com'
```

- 2.5. Verify that the OSD is up.