

```
[ceph: root@clienta /]# ceph osd pool set testpool pg_autoscale_mode off
set pool 6 pg_autoscale_mode to off
[ceph: root@clienta /]# ceph osd pool set testpool pg_num 8
set pool 6 pg_num to 8
[ceph: root@clienta /]# ceph osd pool autoscale-status
```

POOL	SIZE	TARGET SIZE	RATE	RAW CAPACITY	RATIO	TARGET
RATIO	EFFECTIVE RATIO	BIAS	PG_NUM	NEW PG_NUM	AUTOSCALE	
device_health_metrics	0		3.0	92124M	0.0000	
	1.0	1	on			
.rgw.root	1323		3.0	92124M	0.0000	
	1.0	32	on			
default.rgw.log	3702		3.0	92124M	0.0000	
	1.0	32	on			
default.rgw.control	0		3.0	92124M	0.0000	
	1.0	32	on			
default.rgw.meta	0		3.0	92124M	0.0000	
	4.0	8	on			
testpool	0		3.0	92124M	0.0000	
	1.0	8	32	off		

```
[ceph: root@clienta /]# ceph health detail
HEALTH_OK
```

- 1.5. Set the PG autoscale option to warn for the pool `testpool`. Verify that cluster health status is now `WARN`, because the recommended number of PGs is higher than the current number of PGs. It might take several minutes before the cluster shows the health warning message.

```
[ceph: root@clienta /]# ceph osd pool set testpool pg_autoscale_mode warn
set pool 6 pg_autoscale_mode to warn
[ceph: root@clienta /]# ceph health detail
HEALTH_WARN 1 pools have too few placement groups
[WRN] POOL_T00_FEW_PGS: 1 pools have too few placement groups
Pool testpool has 8 placement groups, should have 32
```

- 1.6. Enable the PG autoscale option and verify that the number of PGs has been increased automatically to 32, the recommended value. This increase might take a few minutes to display.

```
[ceph: root@clienta /]# ceph osd pool set testpool pg_autoscale_mode on
set pool 6 pg_autoscale_mode to on
[ceph: root@clienta /]# ceph osd pool autoscale-status
```

POOL	SIZE	TARGET SIZE	RATE	RAW CAPACITY	RATIO	TARGET
RATIO	EFFECTIVE RATIO	BIAS	PG_NUM	NEW PG_NUM	AUTOSCALE	
device_health_metrics	0		3.0	92124M	0.0000	
	1.0	1	on			
.rgw.root	1323		3.0	92124M	0.0000	
	1.0	32	on			
default.rgw.log	3702		3.0	92124M	0.0000	
	1.0	32	on			