- The `serverc`, `serverd`, and `servere` nodes comprise an operational 3-node Ceph cluster. All three nodes operate as a MON, a MGR, and an OSD host with three 10 GB collocated OSDs.

Warning
The parameters used in this exercise are appropriate for this lab environment. In production, these parameters should only be modified by qualified Ceph administrators, or as directed by Red Hat Support.

1. Log in to `clienta` as the `admin` user. Create a new pool called `testpool`, set the PG autoscale mode to `warn`, reduce the number of PGs, and view the health warning messages. Set the PG autoscale mode to `on` again, and then verify the number of PGs and that cluster health is `ok` again.

1.1. Connect to `clienta` as the `admin` user and use `sudo` to run the `cephadm` shell.

```
[student@workstation ~]$ ssh admin@clienta
[admin@clienta ~]$ sudo cephadm shell
[ceph: root@clienta /]#
```

1.2. Create a new pool called `testpool` with the default number of PGs.

```
[ceph: root@clienta /]# ceph osd pool create testpool
pool 'testpool' created
```

1.3. Verify the cluster health status and the information from the PG autoscaler. The autoscaler mode for the created pool `testpool` should be `on` and the number of PGs is 32.

```
[ceph: root@clienta /]# ceph health detail
HEALTH_OK
[ceph: root@clienta /]# ceph osd pool autoscale-status
POOL                    SIZE  TARGET SIZE  RATE  RAW CAPACITY   RATIO  TARGET
 RATIO  EFFECTIVE RATIO  BIAS  PG_NUM  NEW PG_NUM  AUTOSCALE
device_health_metrics      0                 3.0        92124M  0.0000
                         1.0       1              on
.rgw.root                1323                3.0        92124M  0.0000
                         1.0      32              on
default.rgw.log          3702                3.0        92124M  0.0000
                         1.0      32              on
default.rgw.control         0                3.0        92124M  0.0000
                         1.0      32              on
default.rgw.meta            0                3.0        92124M  0.0000
                         4.0       8              on
testpool                    0                3.0        92124M  0.0000
                         1.0      32              on
```

1.4. Set the PG autoscale option to `off` for the pool `testpool`. Reduce the number of PGs to 8. Verify the autoscale recommended number of PGs, which should be 32. Verify that the cluster health is OK.

有版权 禁止外传