```
[ceph: root@clienta /]# ceph tell osd.0 bluestore allocator score block
{
    "fragmentation_rating": 0.0016764709285897418
}
```

2. By default, Red Hat Ceph Storage allows one PG backfill at a time, to or from an OSD. Modify this parameter to 2 on a per-OSD basis. Configure PG backfilling on an OSD.

2.1. Select one OSD running on the `serverc` node and obtain its IDs. In the following example, the options are the `osd.0, osd.1` and `osd.2` OSDs. Yours might be different.

```
[ceph: root@clienta /]# ceph osd tree
ID  CLASS  WEIGHT   TYPE NAME         STATUS  REWEIGHT  PRI-AFF
-1         0.08817  root default
-3         0.02939      host serverc
 0    hdd  0.00980          osd.0         up   1.00000  1.00000
 1    hdd  0.00980          osd.1         up   1.00000  1.00000
 2    hdd  0.00980          osd.2         up   1.00000  1.00000
-7         0.02939      host serverd
 3    hdd  0.00980          osd.3         up   1.00000  1.00000
 5    hdd  0.00980          osd.5         up   1.00000  1.00000
 7    hdd  0.00980          osd.7         up   1.00000  1.00000
-5         0.02939      host servere
 4    hdd  0.00980          osd.4         up   1.00000  1.00000
 6    hdd  0.00980          osd.6         up   1.00000  1.00000
 8    hdd  0.00980          osd.8         up   1.00000  1.00000
```

2.2. On your selected OSD on host `serverc`, retrieve the value for the `osd_max_backfills` parameter. In this example, the selected OSD is `osd.0`.

```
[ceph: root@clienta /]# ceph tell osd.0 config get osd_max_backfills
{
    "osd_max_backfills": "1"
}
```

2.3. Modify the current runtime value for the `osd_max_backfills` parameter to 2.

```
[ceph: root@clienta /]# ceph tell osd.0 config set osd_max_backfills 2
{
    "success": "osd_max_backfills = '2' "
}
```

3. By default, Red Hat Ceph Storage allows three simultaneous recovery operations for HDDs and ten for SSDs. Modify the maximum number of data recovery operations to 1 per OSD.

3.1. Verify the value of the `osd_recovery_max_active` parameter on the OSD of your choice. The default value for the `osd_recovery_max_active` is 0, meaning that the values in `osd_recovery_max_active_hdd` and `osd_recovery_max_active_ssd` are used instead.