

## ► Solution

# Tuning and Troubleshooting Red Hat Ceph Storage

In this lab, you will tune and troubleshoot Red Hat Ceph Storage.

## Outcomes

You should be able to:

- Troubleshoot Ceph clients and Ceph performance issues.
- Modify the logging options for Ceph clients.
- Tune the recovery and backfill processes to preserve cluster performance.

## Before You Begin

As the student user on the workstation machine, use the `lab` command to prepare your system for this lab.

This command ensures that the lab environment is available for the exercise.

```
[student@workstation ~]$ lab start tuning-review
```

## Instructions

1. Log in to `clienta` as the `admin` user. Verify the health of the Ceph storage cluster. Hypothesize possible causes for the displayed issues.

- 1.1. Log in to `clienta` as the `admin` user and use `sudo` to run the `cephadm` shell.

```
[student@workstation ~]$ ssh admin@clienta
[admin@clienta ~]$ sudo cephadm shell
[ceph: root@clienta /]#
```

- 1.2. Verify the status of the storage cluster.

```
[ceph: root@clienta /]# ceph status
cluster:
  id:         2ae6d05a-229a-11ec-925e-52540000fa0c
  health: HEALTH_WARN
           1 failed cephadm daemon(s)
           clock skew detected on mon.serverd
           1 osds down
           Reduced data availability: 4 pgs inactive, 19 pgs peering
           Degraded data redundancy: 35/567 objects degraded (6.173%), 7 pgs
degraded

services:
```