

```

mon: 4 daemons, quorum serverc.lab.example.com,clienta,serverd,servere (age
33s)
mgr: serverc.lab.example.com.aiqepd(active, since 9m), standbys:
clienta.nncugs, servere.kjwyko, serverd.klrkci
osd: 9 osds: 8 up (since 32s), 9 in (since 8m)
rgw: 2 daemons active (2 hosts, 1 zones)

data:
pools: 5 pools, 105 pgs
objects: 189 objects, 4.9 KiB
usage: 105 MiB used, 90 GiB / 90 GiB avail
pgs: 18.095% pgs not active
    35/567 objects degraded (6.173%)
    68 active+clean
    19 peering
    11 active+undersized
    7 active+undersized+degraded

```

1.3. Verify the health details of the storage cluster.

Two separate issues need troubleshooting. The first issue is a clock skew error, and the second issue is a down OSD that is degrading the PGs.

```

[ceph: root@clienta /]# ceph health detail
HEALTH_WARN 1 failed cephadm daemon(s); clock skew detected on mon.serverd; 1
osds down; Degraded data redundancy: 63/567 objects degraded (11.111%), 15 pgs
degraded; 44 slow ops, oldest one blocked for 79 sec, mon.serverd has slow ops
[WRN] CEPHADM_FAILED_DAEMON: 1 failed cephadm daemon(s)
    daemon osd.4 on servere.lab.example.com is in unknown state
[WRN] MON_CLOCK_SKEW: clock skew detected on mon.serverd
    mon.serverd clock skew 299.62s > max 0.05s (latency 0.0155988s)
[WRN] OSD_DOWN: 1 osds down
    osd.4 (root=default,host=servere) is down
[WRN] PG_DEGRADED: Degraded data redundancy: 63/567 objects degraded (11.111%), 15
pgs degraded
    pg 3.1 is active+undersized+degraded, acting [0,7]
...output omitted...
    pg 4.15 is active+undersized+degraded, acting [0,7]

```

2. First, troubleshoot the clock skew issue.

- 2.1. Open a second terminal and log in to serverd as the admin user. The previous health detail output stated that the time on the serverd system is 300 seconds different from the other servers. View the chronyd service status on the serverd system to identify the problem.

The chronyd service is inactive on the serverd system.