

```
[admin@node ~]# ceph auth get-or-create client.glance mon 'profile rbd' \
osd 'profile rbd pool= images' mgr 'profile rbd pool=images'
```

## Image Storage

In OpenStack, the default back end for the Image service is a file store that is located with the Glance API node on the controller node. The location is configurable, with a default of `/var/lib/glance`. To improve scalability, the Image service implemented an image cache at the default `/var/lib/glance/image-cache/` location on the controller node. When the Compute service loads images that are stored in the default QCOW2 format to convert to RAW for use on the compute nodes, then the converted image is cached.

When Red Hat OpenStack Platform is installed with the Swift Object Store, TripleO places the image service back end on Swift by default. The Swift service creates a container called `glance` for storing Glance images.

When Ceph storage is integrated into RHOSP, TripleO places the image service back end on Ceph RADOS Block Devices (RBD) by default. Glance images are stored in a Ceph pool called `images`. RHOSP works with images as immutable blobs and handles them accordingly. The pool name is configurable with the `glance_pool_name` property. The `images` pool is configured as a replicated pool by default, which means that all images are replicated across storage devices for transparent resilience.

An image pool can be configured as erasure-coded to conserve disk space with a slight increase in CPU utilization.

When using Ceph as the storage back end, it is important to disable the image cache, as it is not needed because Ceph expects Glance images to be stored in the RAW format. When using RAW images, all image interactions occur within Ceph, including image clone and snapshot creation. Disabling the image cache eliminates significant CPU and network activity on controller nodes.

When using a distributed architecture with Distributed Compute Nodes (DCN), TripleO can configure the Image service with an image pool at each remote site. You can copy images between the central (hub) site and the remote sites. The DCN Ceph cluster uses RBD technologies, such as copy-on-write and snapshot layering, for fast instance launching. The Image, Block Storage, and Compute services must all be configured to use Ceph RBD as their back-end storage.

## Object Storage

Object storage is implemented in OpenStack by the Object Store service (Swift). The Object Store service implements both the Swift API and the Amazon S3 API. The default storage back end is file-based, and uses an XFS-formatted partition mount in subdirectories of `/srv/node` on the designated storage node. You can also configure the Object Store service to use an existing, external Swift cluster as a back end.

When Ceph storage is integrated into RHOSP, TripleO configures the Object Store service to use the RADOS Gateway (RGW) as the back end. Similarly, the Image service is configured for RGW because Swift would not be available as a back end.

The Ceph Object Gateway can be integrated with the Keystone identity service. This integration configures RGW to use the Identity service as the user authority. If Keystone authorizes a user to access the gateway, then the user is also initially created on the Ceph Object Gateway. Identity tokens that Keystone validates are considered valid by the Ceph Object Gateway. The Ceph Object Gateway is also configured as an object-storage endpoint in Keystone.