```
[ceph: root@clienta /]# ceph osd pool create cephfs.cl260-fs.data
pool 'cephfs.cl260-fs.data' created
[ceph: root@clienta /]# ceph osd pool create cephfs.cl260-fs.meta
pool 'cephfs.cl260-fs.meta' created
[ceph: root@clienta /]# ceph osd pool ls
...output omitted...
cephfs.cl260-fs.data
cephfs.cl260-fs.meta
```

1.3. Create a CephFS file system called `cl260-fs`.

```
[ceph: root@clienta /]# ceph fs new cl260-fs cephfs.cl260-fs.data \
cephfs.cl260-fs.meta
new fs with metadata pool 14 and data pool 15
[ceph: root@clienta /]# ceph fs ls
name: cl260-fs, metadata pool: cephfs.cl260-fs.data, data pools: [cephfs.cl260-
fs.meta ]
```

1.4. Create an MDS service called `cl260-fs` with an MDS instance on `serverc`.

```
[ceph: root@clienta /]# ceph orch apply mds cl260-fs \
  --placement="2 serverc.lab.example.com serverd.lab.example.com"
Scheduled mds.cl260-fs update...
[ceph: root@clienta /]# ceph orch ps --daemon-type mds
NAME                         HOST                  STATUS         REFRESHED ...
mds.cl260-fs.serverc.iuwwzt  serverc.lab.example.com  running (53s)  46s ago   ...
mds.cl260-fs.serverd.lapeyj  serverd.lab.example.com  running (50s)  47s ago   ...
```

1.5. Verify that the MDS service is up and running.

```
[ceph: root@clienta /]# ceph mds stat
cl260-fs:1 {0=cl260-fs.serverd.lapeyj=up:active} 1 up:standby
[ceph: root@clienta /]# ceph status
  cluster:
    id:     2ae6d05a-229a-11ec-925e-52540000fa0c
    health: HEALTH_OK

  services:
...output omitted...
    mds: 1/1 daemons up, 1 standby
...output omitted...
```

2. Install the `ceph-common` package. Mount the CephFS file system to the `/mnt/cephfs` directory on the `clienta` host. Save the key-ring associated with the `client.admin` user to the `/root/secretfile` file. Use this file to authenticate the mount operation. Verify that the `/mnt/cephfs` directory is owned by the `admin` user.

2.1. Exit the `cephadm` shell and switch to the `root` user. Install the `ceph-common` package.